

Scope & Audience

This guideline is intended for developers, AI engineers, DevOps engineers, architects, and AI coding assistants working on GCP-based projects. It prescribes Infrastructure as Code (IaC) and Continuous Integration / Continuous Deployment (CI/CD) best practices to enable code components to autonomously provision their required cloud infrastructure.

1. Layered Architecture

Define clear separation of concerns across five layers. Each layer consumes services from the layer beneath and exposes well-documented interfaces to layers above.

1.1 Pipeline Layer

- **Data / Knowledge Base Pipeline:** Automated ingestion, transformation, and versioning of data.
- **Code / Deployment Pipeline:** IaC templates (Terraform, Deployment Manager) stored in Git, triggered by commits.
- **Cloud Operations Pipeline:** Monitoring, logging, alerting, and automated rollbacks via Cloud Build and Cloud Functions.

1.2 Tools Layer

- **Internal Tools:** Custom-made, GCP tenant-scoped service catalog entries.
- **External Tools:** Third-party APIs/MCPs, managed connectors.

1.3 Models Layer

- Packaging and versioning of AI/ML models.
- Artifact storage in Container Registry / Artifact Registry.

1.4 Agents Layer

- Deployment of agent runtimes (Cloud Run, Cloud Functions).
- Definition of agent permissions via IAM roles and service accounts.

1.5 Workflows Layer

- Orchestration via Workflows, Composer, or Cloud Tasks.
- YAML-driven workflow definitions stored alongside code.

2. Simplified Environment Transitions

Minimize friction when moving code and IaC between dev, test, and pre-prod.

1. Parameterized Configuration

- Store environment-specific values (project IDs, region, service account emails) in Terraform variables or Cloud Build substitution maps.
- Use separate tfvars files: dev.tfvars, test.tfvars, preprod.tfvars.

2. Service Accounts & API Keys

- Define one service account per environment; rotate keys via Secret Manager.
- Automate Secret Manager replication across environments using Terraform.
- Reference secrets in pipelines with gcloud secrets versions access latest in Cloud Build.

3. GitOps Branching Strategy

- main → pre-prod
- develop → test
- feature/* → dev
- Merge rules trigger automated environments via Cloud Build triggers.

4. Shared Modules & Blueprints

- Publish reusable Terraform modules in a private Artifact Registry.
- Enforce module versioning to ensure stability across environments.

3. IaC & CI/CD Principles

Enable code-driven infrastructure provisioning with automated quality gates.

1. Declarative IaC

- Use Terraform (recommended) or Deployment Manager for all infra.
- Enforce style and policy with tflint, terraform fmt, and Conftest (OPA).

2. Pipeline Automation

- Cloud Build pipelines defined in cloudbuild.yaml.
- Steps: terraform init → terraform plan → policy checks → manual approval (test/preprod) → terraform apply.

3. Immutable Infrastructure

- Avoid in-place resource mutation; use resource replacement and canary deployments.
- For compute workloads, use container images with versioned tags.

4. Automated Testing

- Unit tests for Terraform modules via terraform-compliance.
- Integration tests using Test Kitchen or Molecule against deployed stacks.

5. Rollbacks and Roll-forwards

- Leverage native Terraform state and versioning.
- Cloud Build rollback on failed smoke tests; alerting via Cloud Monitoring.

4. Maintainability & Continuous Improvement

Focus on simplicity, reusability, and evolutionary enhancements.

- **Minimal Deployments:** Combine related resources in modules to reduce deployment time.
- **Layer Reusability:** Publish shared modules for networking, IAM, storage, and AI services.
- **Strategic Use of Layers:** Keep pipelines, tools, models, agents, and workflows decoupled; evolve each layer independently.
- **Documentation as Code:** Maintain pipeline, module, and API docs in Markdown alongside code.
- **Versioning & Changelog:** Tag every release; publish changelogs automatically via CI/CD pipelines.

5. Compliance, Budget Management & Scalability

Ensure governance, cost control, and the ability to serve up to 10,000 users.

5.1 Compliance Controls

- Enforce organization policies with Org Policy Service (e.g., allowed regions, machine types).
- Enable Audit Logging on all projects.
- Periodic configuration scans via Forseti / Config Connector.

5.2 Budget Management

- Define budgets and alerts in Billing Budget API.
- Set cost ceilings per environment; automate notifications via Cloud Functions.
- Leverage Quota Management to cap API usage.

5.3 Scalability

- Compute: Use Cloud Run or GKE with Horizontal Pod Autoscaler.
 - Storage: Choose multi-regional Cloud Storage with lifecycle rules.
 - Database: Use managed services (Cloud SQL with read replicas, Bigtable for high throughput).
 - Networking: Auto-scale load balancers with global backend services.
-

6. Actionable Recommendations & Best Practices

- **Adopt GitOps:** Treat all infra changes as pull requests with peer reviews.
- **Enforce Policy as Code:** Embed security/compliance rules in CI using OPA/Gatekeeper.
- **Automate Secrets Management:** Rotate service credentials automatically; never store creds in code.
- **Use Feature Toggles:** Deploy infra early; gate new services via runtime flags.
- **Regularly Review Costs:** Schedule automated budget reports; optimize idle resources.
- **Monitor & Alert:** Define SLOs/SLIs; set up dashboards and alerts for key metrics.
- **Perform Chaos Testing:** Use Fault Injection frameworks to validate resilience.
- **Continuous Learning:** Hold quarterly IaC and pipeline reviews; adopt community modules and patterns.

